



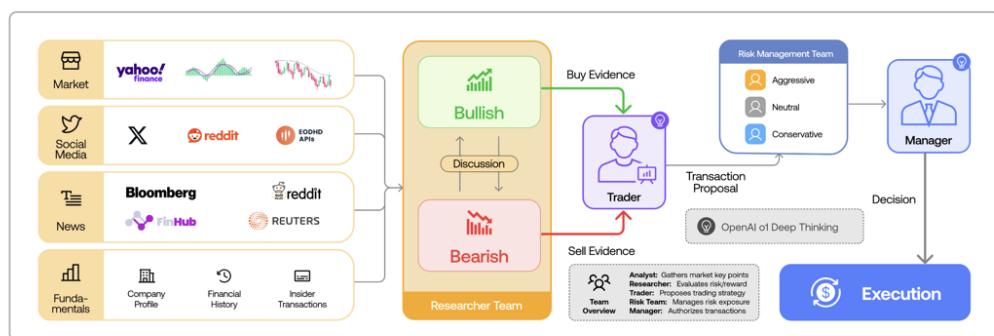
# AI Investor Systems: Multi-Agent Architectures and Performance Review

## Introduction

Artificial intelligence is rapidly transforming investment management, enabling **AI investor systems** that autonomously analyze markets, execute trades, and manage risk with minimal human input. These systems combine advanced machine learning (e.g. deep reinforcement learning) with structured multi-agent designs to emulate a team of expert traders. The goal is to outperform traditional benchmarks (such as the S&P 500 index) by a wide margin while tightly controlling risk. Recent developments in **persona-based agents**, multi-asset trading algorithms, and large language model (LLM)-integrated pipelines have brought this vision closer to reality. In this report, we review how deep-thinking AI investor systems are being developed in academia and industry, focusing on their architectures, asset scope, and reported performance. We compare academic prototypes to real-world implementations (hedge funds, fintech startups, open-source projects) to understand the design choices and efficacy of these systems – ultimately informing how one might build a personal AI investor on their own infrastructure.

## Persona-Based and Modular Agent Architectures

A prevailing design pattern in cutting-edge trading AI is the use of **multiple specialized agents or “personas”** that mimic the division of labor on a trading desk. Instead of a monolithic black-box model, these systems break the problem into modular roles – for example, separate agents for market analysis, strategy formulation, risk management, and execution [1](#) [2](#). Each agent acts as an expert in its domain, and together they collaborate (often via structured communication or debates) to reach trading decisions. This approach improves robustness by enabling cross-validation and checks and balances between agents [3](#). It also lends itself to greater transparency, since each agent’s contribution can be interpreted.



*Example of a multi-agent “deep thinking” trading architecture, where specialized agents (Analysts, Bullish vs. Bearish Researchers, Trader, Risk Managers, etc.) debate and collaborate to propose and execute trades. This persona-based design mirrors human investment teams and provides modularity, explainability, and improved decision quality.* [4](#) [5](#)

**Persona-based systems** often include agents modeled after different trading styles or perspectives. For instance, a **day-trader agent** might focus on high-frequency technical signals, while a **swing-trader agent** analyzes multi-day trends – each providing recommendations that a higher-level arbiter can reconcile. A recent conceptual demo by Khan (2025) showcased a “deep thinking” trading workflow with distinct sub-agents assuming roles such as: *Analysts* (gathering 360° market data), *Researcher* agents labeled “Bull” and “Bear” who debate the outlook, a *Trader* agent translating the consensus into a concrete strategy, and a *Risk Management team* with personas like “Risky,” “Safe,” and “Neutral” providing multi-angle risk checks <sup>6</sup> <sup>7</sup>. Finally, a manager agent approves the trade, ensuring it aligns with overall objectives before execution. By dividing responsibilities in this manner, the system can emulate the deliberative process of a human trading committee, which has been argued to produce more balanced and resilient decisions than any single model <sup>8</sup> <sup>9</sup>.

Crucially, multi-agent architectures require a coordination mechanism for agents to share information and resolve disagreements. Some frameworks use an explicit **controller or facilitator agent** to orchestrate the workflow (assign tasks, integrate outputs, etc.), while others rely on iterative dialogues between agents (often leveraging LLMs to enable natural language “debates”). Research indicates that when agents with different expertise collaborate (e.g. technical analysis, fundamental analysis, news sentiment), the ensemble can outperform both single-agent models and even market index benchmarks <sup>10</sup>. For example, an explainable multi-modal crypto trading AI with specialized sub-agents was shown to **beat standard market indices on multiple metrics** by combining each agent’s insights <sup>10</sup>. This evidence supports the intuition that a team of focused “persona” agents, if properly coordinated, yields a more **adaptive and fault-tolerant** trading system than a lone generalist agent <sup>11</sup> <sup>13</sup>.

## Multi-Asset Class Coverage (Equities, Crypto, and Options)

AI investor systems are being applied across diverse asset classes – from traditional equities to cryptocurrencies and complex derivatives – often within the same unified framework. Supporting **multiple asset classes** is challenging because each market has distinct dynamics (e.g. equities respond to earnings and news, crypto to on-chain metrics and sentiment, options to volatility and time-decay). Modern AI agents address this by either training specialized models per asset class or designing architectures that can ingest multi-modal data covering all relevant domains.

**Equities:** Academic projects like P1GPT (2025) focus on U.S. stock trading by fusing technical price patterns, company fundamentals, and real-time news analysis through a team of LLM-based agents <sup>12</sup> <sup>13</sup>. P1GPT’s layered multi-agent pipeline included dedicated agents for **fundamental analysis** (processing financial statements, valuations), **technical analysis** (market signals like momentum or RSI), and **news sentiment** – all coordinated to produce stock buy/sell decisions with explanations <sup>14</sup> <sup>15</sup>. This multi-modal approach proved effective in backtests on equities, as described later in the performance section. Similarly, the open-source **FinRL** framework provides a configurable pipeline for stock trading using deep reinforcement learning, and it has been extended to include **LLM-generated signals** (e.g. sentiment from news or social media) as additional inputs for trading agents <sup>16</sup>. This integration of text-based intelligence aims to give equity trading agents a broader awareness of market-moving information beyond price data.

**Cryptocurrency:** Crypto markets are known for high volatility and rich alternative data (on-chain transactions, decentralized finance (DeFi) metrics, social media hype). AI systems have embraced these by using multi-agent designs that handle both *off-chain* and *on-chain* data. For example, a **crypto portfolio RL agent** (CryptoRLPM) incorporated blockchain analytics (whale wallet movements, network activity)

alongside price history <sup>17</sup>. By structuring its pipeline into modules (data feed, refinement, portfolio optimization, etc.), it dynamically adjusted a crypto portfolio and achieved **83% higher cumulative returns than a Bitcoin-only benchmark**, with significantly better risk-adjusted returns (Sortino ratio) <sup>18</sup>. This underscores the benefit of cross-domain signals in crypto – an AI that “sees” on-chain liquidity shifts or protocol metrics can anticipate moves that price-only models miss. Another study introduced a **multi-agent RL simulator for crypto markets** to train agents that generalize across different price trend scenarios <sup>19</sup>. On the DeFi front, the *SuperIntent* platform (by XY Finance) demonstrates an intent-driven multi-agent system that spans yield farming, lending, and cross-chain trading. Its architecture has dedicated agents for data collection across dozens of blockchains, strategy formulation (e.g. picking the best yield opportunities), and execution on-chain <sup>20</sup> <sup>21</sup>. By handling **cross-chain data and transactions**, such systems aim to optimize crypto investments in a holistic way – something a human would struggle with given the fragmentation of DeFi data and venues <sup>22</sup> <sup>23</sup>.

**Options and Derivatives:** Options trading introduces another layer of complexity due to the need to manage volatility and hedge risks. Cutting-edge research is extending AI agents into this domain as well. *OPHR (Option Portfolio Hedging RL)*, proposed in 2024, is **the first multi-agent RL framework for volatility trading via options** <sup>24</sup>. It employs two cooperating agents – an **Option Position Agent** that decides when to be long or short volatility (i.e. buying or selling options), and a **Hedger Agent** that dynamically selects appropriate hedging strategies to manage risk <sup>25</sup> <sup>26</sup>. By training these agents on historical options data (for Bitcoin and Ethereum options markets), the system learned to time volatility spikes (buying options before big moves, selling premium during calm periods) better than traditional rule-based strategies <sup>27</sup>. Crucially, the multi-agent design allowed separate focus on profit (the position agent maximizing returns) and safety (the hedger minimizing tail risk), striking a balance. Empirical results showed this AI could **outperform benchmark strategies** for both rising and falling volatility, while effectively controlling tail-risk when selling options <sup>28</sup>. This suggests AI agents can handle options and other derivatives by decomposing the tasks (signal generation vs. risk mitigation), and such approaches could be expanded to futures, credit derivatives, and beyond.

The table below summarizes some representative AI trading systems across asset classes, highlighting their approaches and reported performance:

System (Year)	Assets Covered	Approach & Architecture	Reported Performance
<b>P1GPT (2025)</b> <sup>29</sup> <sup>30</sup>	US Equities (multi-modal data)	Multi-agent LLM pipeline (technical, fundamental, news analysts; planning & integration layers)	Annualized returns exceeded baselines by ~20% points; Sharpe ~2.3–3.4; max drawdown ~2–7% (vs. <1.5% for ultra-conservative strategies)
<b>CryptoRLPM (2024)</b> <sup>18</sup>	Crypto Portfolio (BTC, etc.)	Deep RL agent with structured data pipeline (on-chain signals + price); modular training units	+83% cumulative return over BTC benchmark; higher Sortino (risk-adjusted return) indicating superior performance

System (Year)	Assets Covered	Approach & Architecture	Reported Performance
MAMA RL (2025) <small>31</small>	Multi-Asset (Stocks, Bonds, Commodities)	<i>Multi-Asset Multi-Agent RL:</i> Graph neural nets for each asset class + hierarchical asset allocator; tackles concept drift and interest rates	40.9% compounded annual growth rate (CAGR) in backtest, outperforming the next-best method by 23 percentage points. (Significantly above typical 10% market returns)
OPHR (2024) <small>24 28</small>	Options (Crypto options on BTC/ ETH)	Two-agent deep RL (Volatility trader + Hedging agent); collaboration via alternate training; focuses on timing volatility & risk management	Outperformed rule-based vol trading strategies; improved tail- risk management when selling options, indicating safer profit generation during volatile swings
TradingAgents (2025) <small>32 5</small>	US Equities (simulation)	Multi-agent LLM framework emulating a trading firm (teams of analysts for technical, sentiment, news, fundamentals; trader agents; risk agents using debate via LLM prompts)	In a 3-month 2024 simulation, beat baseline strategies (Buy- and-Hold, MACD, RSI, etc.) in <b>cumulative returns and</b> <b>Sharpe ratio</b> , with lower max drawdown, demonstrating superior risk-adjusted performance over single-agent approaches.

Table 1: Examples of AI Investor Systems across asset classes, with architecture highlights and performance (in backtesting) relative to benchmarks.

## Reinforcement Learning and Deep Learning Techniques

Many advanced AI investor systems are underpinned by **deep reinforcement learning (DRL)**, which frames trading as a sequential decision-making problem. In a DRL setup, an agent learns an optimal trading policy by interacting with a simulated market environment: observing state features, executing buy/sell/hold actions, and receiving rewards (often based on portfolio returns or risk-adjusted returns). Over time, the agent learns which actions yield the highest long-term reward. This approach has been applied to portfolio allocation, single-asset trading, order execution, and more 33.

Key RL techniques in this domain include:

- **Policy Gradient and Actor-Critic methods:** Algorithms like Proximal Policy Optimization (PPO) and Deep Deterministic Policy Gradient (DDPG) are popular for trading tasks 34. They can handle continuous action spaces (e.g. allocating fractions of capital among assets) and are well-suited for **portfolio rebalancing problems**. For example, an RL agent might continuously adjust weights of stocks vs. bonds to maximize return for a given risk – tasks which traditional mean-variance methods handle sub-optimally under non-stationary conditions. DRL agents have shown the ability to dynamically **adapt to regime shifts** (bull vs. bear markets) by retraining or via on-line learning, whereas static allocators cannot 35. - **Deep Q-Networks (DQN) and Variants:** Value-based methods (DQN, Double DQN, etc.) have been used for discrete trading actions like deciding to buy, sell, or hold a single asset. These were among the earliest DRL trading

agents (e.g. JPMorgan's LOXM for execution reportedly used DQN variants). However, value-based RL has to carefully handle the **reward design** to avoid “reward hacking.” If the reward is simply profit, an unconstrained agent might learn extreme leverage or overfit to training data. Recent studies incorporate **penalties for risk** (drawdowns or volatility) into the reward function, or use **multi-objective RL** to balance return and risk. A *behaviorally-informed DRL* approach in 2026 explicitly modeled investor loss-aversion and overconfidence in the reward shaping, to guide the agent toward more realistic risk-taking <sup>36</sup> <sup>37</sup>. - **Multi-Agent RL (MARL)**: When dealing with multiple assets or objectives, MARL allows separate agents to learn different facets and then work as a team. The MAMA framework (2025) described earlier is one example, where *intra-asset agents* specialize in picking stocks within a sector and *inter-asset agents* decide allocations between asset classes <sup>38</sup> <sup>39</sup>. They coordinate to form a complete portfolio. MARL can also simulate market dynamics by training agents that trade **against each other** or against synthetic market-makers. For instance, multi-agent simulations have been used to generate more robust trading strategies that are not overfit to one market scenario <sup>40</sup>. Another research direction is using *adversarial agents* to stress-test the main trading agent (one agent tries to maximize profit while an adversary perturbs prices or adds transaction costs), improving the main agent’s robustness <sup>41</sup>. - **Deep Neural Networks for Feature Modeling**: Whether using RL or supervised learning, almost all AI trading systems leverage deep networks to encode market state. Recurrent neural networks or LSTMs capture temporal dependencies in price series; convolutional networks or transformers may be used to detect patterns in time-series or even treat candlestick charts as images. More exotic architectures like graph neural networks have been employed for **relations between assets** (modeling a stock network or crypto token graph) – for example, MAMA used GNNs to learn relationships within each asset class (stocks or bonds) for better diversification <sup>39</sup>. The state representation is critical: it often includes technical indicators, macro data, and sentiment scores. In cutting-edge systems, **LLMs are now generating some of these state features** – e.g. parsing news headlines into sentiment metrics or summarizing SEC filings, which then feed into the numerical model <sup>16</sup>.

One of the strengths of RL-based investors is the ability to directly optimize objectives like **Sharpe ratio or drawdown constraints**. Researchers have experimented with custom reward functions to penalize excess volatility or large peak-to-trough losses. By tuning the reward, an RL agent can be trained to, say, “achieve at least X return while keeping max drawdown under 20%.” However, this is non-trivial: overly constraining risk can simply lead the agent to stay in cash. A practical compromise is to use risk-adjusted rewards (like Sharpe or Sortino ratio as the reward signal), which implicitly reward return and penalize volatility. Many DRL trading implementations monitor metrics like maximum drawdown during training and include them in early stopping criteria or as part of the fitness function <sup>32</sup> <sup>4</sup>. The result is that some learned strategies exhibit much lower drawdowns than naive strategies for comparable returns. For example, P1GPT’s multi-agent system maintained drawdowns below 7% in stocks like GOOGL and TSLA, whereas simpler strategies often saw larger swings (or conversely, lower returns for similar drawdown) <sup>30</sup> <sup>42</sup>.

In summary, deep RL provides the “brain” for many AI investor systems, enabling continuous learning from interaction with markets. Coupled with deep neural networks for perception (feature extraction) and sometimes guided by financial domain knowledge (through reward shaping or constraints), RL agents have demonstrated an ability to *discover non-obvious trading policies*. Still, challenges remain in training stability (financial time-series are noisy and non-stationary) and avoiding overfitting. That’s why researchers emphasize rigorous **benchmarking and backtesting protocols** – e.g. the FinRL community’s contests and libraries to standardize evaluation across stock trading, crypto trading, etc., with proper train-test splits and walk-forward analysis <sup>43</sup> <sup>33</sup>. An AI investor is only as good as its performance on *unseen* market conditions, so robust validation (and in live deployment, risk limits) are essential parts of the design.

## LLM-Integrated Trading Pipelines

The integration of **large language models (LLMs)** into trading systems is a novel trend that merges quantitative algorithms with qualitative reasoning. LLMs (like GPT-4 and BloombergGPT) excel at understanding natural language and generating human-like analysis, which opens up new possibilities for financial AI: they can read news articles, interpret earnings call transcripts, gauge sentiment from social media, and even explain the rationale behind trades. In AI investor architectures, LLMs are being deployed both as **data analysis agents** and as **decision-making facilitators**.

**LLMs as Market Analysts:** Financial LLMs fine-tuned on economic news and reports can serve as additional “analyst agents” in a multi-agent setup. For instance, in the TradingAgents framework, some agents are effectively LLM-powered analysts for specific data streams – e.g. a *News Analyst* agent that uses an LLM to summarize and classify news, or a *Sentiment Analyst* parsing Reddit or Twitter for market mood <sup>44</sup>. BloombergGPT (2023) demonstrated that an LLM trained on financial text can answer questions and classify sentiment with high accuracy <sup>45</sup>. Such models can convert unstructured text into structured signals (e.g. an earnings report into a bullish or bearish score) that feed into trading decisions. In P1GPT’s architecture, the Analysis Layer included *Intelligent Specialized Agents (ISAs)* like a **News ISA** and a **Fundamental ISA**, which presumably use LLM capabilities to interpret qualitative information (news events, fundamental data) and produce a summary or recommendation <sup>14</sup> <sup>46</sup>. The outputs of these LLM-driven agents are then fused with technical signals in the integration layer. Notably, P1GPT found that this structured fusion of LLM-derived insights with numeric data led to **better returns and lower drawdowns** than purely quantitative baselines <sup>12</sup> <sup>47</sup>, highlighting the value of incorporating rich context that LLMs can provide.

**Role-playing and Reasoning Chains:** Some systems explicitly use LLMs to *simulate a discussion* between different expert personas. A prompt might be engineered where the LLM is instructed to adopt the role of a cautious investor and an aggressive investor debating a particular trade, with the model outputting a conclusion. While earlier attempts in 2024–2025 used this “role-play” method informally, more recent work has formalized it. For example, *TradingGPT* (not to be confused with *TradingAgents*) proposed a multi-agent LLM system with **distinct characters** that debate and reason, combined with a shared memory for financial context <sup>48</sup>. The **ReAct prompting framework** is often used, where agents first reason (in natural language) about their observations, then take tool-based actions (like querying a pricing API), then reason again, and so on <sup>49</sup>. This chain-of-thought approach allows complex reasoning like: “The Fed hinted at rate hikes – the macro agent predicts market downturn; the technical agent sees a breakdown of support levels – collectively suggesting we sell.” The *TradingAgents* researchers noted that by using LLMs with ReAct, the system enabled explicit reasoning and justifications for trades, which is valuable for explainability <sup>50</sup> <sup>4</sup> <sup>6</sup>. In their tests, the natural language dialogue between agents (and a final “reflective” agent overseeing the discussion) produced human-readable explanations for each trade, satisfying the need for **transparency** in an otherwise black-box AI <sup>51</sup>.

**Pipeline Orchestration:** LLMs can also act as high-level orchestrators in a trading pipeline. One design is to have the LLM analyze the outputs of various narrow models and then decide a course of action or generate a plan. The P1GPT system exemplified this by dividing into layers: a *Planning Layer* (possibly LLM-driven) that takes the raw intent or goal (e.g. “maximize returns with moderate risk”) and delegates tasks to specialized agents <sup>15</sup>. After analysis, an *Integration/Reasoning Layer* (again LLM-coordinated) compiles the agents’ findings into a coherent strategy recommendation <sup>15</sup> <sup>46</sup>. The final Decision Layer then outputs the trade action along with an *LLM-generated explanation* <sup>46</sup>. This structured workflow contrasts with prompting an LLM to “act as a trading bot” in one go – instead, it uses LLMs in a **toolified, modular fashion**, which the

authors argue is more scalable and reliable <sup>52</sup> <sup>53</sup>. Results from P1GPT's backtest on stocks showed the approach yielded **smoother equity curves and consistent risk-adjusted returns**, owing to the system's ability to interpret multi-modal signals and avoid blindly following any single indicator <sup>54</sup> <sup>42</sup>. In short, the LLM served as the glue holding together various AI components, ensuring they reason in concert.

**Real-world LLM usage:** In practice, some fintech trading platforms have begun incorporating LLMs to enhance user experience and strategy discovery. Examples include AI "copilots" that can answer a trader's questions ("Why did the AI sell Apple yesterday?") or even accept high-level instructions ("Invest for me but avoid high-volatility stocks") and translate them into portfolio adjustments – akin to *intent-based investing* <sup>55</sup> <sup>56</sup>. The SuperIntent DeFi system mentioned earlier is built around *user intents*: a user states a goal ("maximize my yield with low risk"), and the AI agents figure out the actions <sup>55</sup> <sup>23</sup>. LLMs are adept at interpreting such natural language intents and mapping them to a set of financial actions (e.g. allocate X% to a stablecoin lending protocol). They can also generate explanations of what they're doing in plain English, which increases trust. On the backend, however, execution is handed off to the specialized agents or hard-coded logic, since LLMs themselves do not interface with APIs unless specifically integrated. Projects like **FinGPT** and **ChatGPT Finance** (community-driven efforts) are exploring using GPT-4 to parse **SEC filings, news headlines, and social media**, providing a higher level summary sentiment that can feed into trading signals <sup>57</sup> <sup>16</sup>.

One challenge with LLM integration is **latency and cost** – running a GPT-4 model for every tick or every trade decision can be slow and expensive. Therefore, current designs often have LLM agents operating on a slower timescale (e.g. analyzing daily news or providing a morning briefing of market risks), complementing faster numerical models that react to intraday price movements. This aligns with a **hands-off, low-frequency approach** – many AI investor systems aim to require only occasional human oversight (checking in once or twice a day). They typically operate on medium-term horizons (hours to days for trades) rather than high-frequency scalping, which would be infeasible for an LLM in the loop <sup>58</sup>. The advantage is that for *longer-term and high-level decision making*, LLMs can incorporate rich context that purely technical models might ignore. As the technology improves (with smaller, faster models and more efficient prompting), we can expect more seamless blending of LLM-driven insight with algorithmic execution.

## Achieving High Returns with Controlled Risk (Performance vs. SPY)

A primary measure of success for an AI investor is whether it can **beat standard benchmarks** like the S&P 500 (often proxied by the SPY ETF) while maintaining reasonable risk. The user's target of "2× annualized returns of SPY with max drawdown <20%" is an ambitious bar – roughly, achieving ~20%+ yearly returns (since the S&P historically returns ~10% annually) and never losing more than 20% from peak. Both academic results and some real-world AI funds suggest this level of performance is *potentially attainable*, at least in backtests and certain market regimes, though not guaranteed.

From the academic side, several prototypes claim **high double-digit annual returns with limited drawdowns**:

- P1GPT's multi-agent LLM system, for example, delivered annualized returns of **25–52%** on test assets (AAPL, GOOGL, TSLA) over an 8-month 2025 period <sup>29</sup>. Its max drawdowns were impressively low (2% on Apple, <7% on Google and Tesla) <sup>30</sup> – far below the S&P 500's typical drawdowns in volatile times. This was achieved through adaptive trading that avoided major downturns and re-entered after dips <sup>59</sup> <sup>60</sup>. The Sharpe ratios above 2.0 indicate strong risk-adjusted performance <sup>61</sup>. While

this is just a backtest, it shows the *potential* for a well-designed AI to double or triple the index returns **without excessive risk**, by actively managing positions and using diverse information. Notably, P1GPT's trades were not very high frequency – it favored a “rhythm of position building and unwinding” with risk-aware exits, which helped in **drawdown control** <sup>59</sup> <sup>62</sup>.

- The MAMA multi-asset RL framework similarly reported ~**40.9% CAGR** in a portfolio of stocks, bonds, and commodities <sup>31</sup>. Details on its drawdown weren't provided in the snippet, but given it surpassed other methods by 23% and presumably had risk controls, we can infer it wasn't achieved by simply leveraging up (the mention of addressing concept drift and including risk-free rate suggests a focus on stable improvements). A 40% CAGR over multiple years would indeed dwarf the S&P 500 (unless one cherry-picks a bull run), so if robust, this is evidence that **multi-agent RL can exploit cross-asset opportunities** to greatly outperform a static index.
- In the crypto domain, doubling the S&P's returns is less of a stretch (since crypto can have triple-digit yearly swings). The CryptoRLPM study's +83% cumulative gain over Bitcoin was in a relatively short horizon <sup>18</sup>. And a **volatility trading RL** might achieve outsized returns by capturing tail events. However, the key is the drawdown: AI systems strive to **limit losses** via built-in risk management agents and conservative mode switches. For instance, OPHR's focus on hedging meant it handled market crashes better than naive strategies (selling options can blow up in crises, but the hedger agent prevents that) <sup>28</sup> <sup>63</sup>.

Real-world evidence is a bit more guarded (since live performance can diverge from backtest). Nonetheless, there are AI-driven funds and portfolios that have claimed strong results. According to Eurekahedge data, **hedge funds using AI** for investment decisions *outperformed the industry average* in most years of the last decade (except 2012) <sup>64</sup>. Some established AI-managed funds delivered **stable double-digit annual returns with strict downside protection**, though others have been volatile <sup>65</sup>. For example, the Hong Kong-based fund Aidyia (run by an AI with minimal human input) was reported to have averaged about **25% yearly returns** in historical testing <sup>66</sup> – roughly 2x the long-term market rate – and with heavy emphasis on avoiding large losses. Aidyia's AI traded fully autonomously and its creators were confident in its ability to adapt to changing market conditions without human intervention <sup>67</sup> <sup>68</sup>. While actual live track records are often proprietary, the presence of AI funds like those suggests that the 2x SPY, <20% drawdown target is not pure fantasy. In fact, analysts noted some AI funds achieved those results by being *very disciplined in downside risk* (cutting positions to cash in bad times) and taking advantage of many uncorrelated signals <sup>65</sup>.

It's important to stress that **backtest success doesn't guarantee future performance** – markets evolve, and an AI can fail if regime changes outpace its learning. However, the modular AI investor systems being developed have a few advantages in sustaining performance: (1) They continuously learn and update (e.g. reinforcement learning agents retrain on new data, or new data feeds get integrated), so they're not fixed to a historical pattern <sup>69</sup>. (2) Their multi-agent structure gives a form of *risk management by design* – e.g. a dedicated risk agent can veto a trader agent's decision if it violates the drawdown or exposure limits <sup>70</sup> <sup>71</sup>. For instance, in the TradingAgents framework, before any trade is executed the **Risk Management Agent** checks the portfolio impact and will reject trades that breach volatility or correlation thresholds <sup>71</sup>. This kind of rule can directly enforce a “max 20% drawdown” constraint by preventing the portfolio from becoming too imbalanced or leveraged. (3) Many systems use **ensemble and hedging techniques**: if one strategy agent is bullish and another is bearish, the disagreement itself can be a signal to reduce position size. This internal hedging can smooth out performance and avoid big directional bets unless there's high consensus <sup>72</sup> <sup>73</sup>.

In summary, the combination of high returns and controlled risk is exactly what these AI systems are engineered to deliver. Academic results are promising, with several multi-agent AI prototypes handily beating market indices in simulation while keeping risk metrics in check. Real-world deployed AI funds have seen periods of exceptional performance as well, though with some cautionary tales (volatile returns for some, and the constant need to validate models in live markets) <sup>65</sup>. The **consensus in the industry** is that AI-driven strategies can indeed “**outperform human traders**” by reacting faster, finding subtle inefficiencies, and removing emotion <sup>74</sup>. They also excel at **automated risk management**, continuously rebalancing to reduce exposure when volatility picks up <sup>75</sup>. The ultimate goal of doubling the market with limited drawdown essentially means achieving a high **Return-over-Maximum-Drawdown (RoMaD)** ratio. Traditional funds strive for  $\text{RoMaD} > 1$  (meaning returns exceed drawdown). AI systems like the ones discussed often report RoMaDs well above 1 – for example, P1GPT on TSLA had ~36% cumulative return vs ~7% max drawdown, a ratio ~5 <sup>54</sup> <sup>30</sup>. Such figures are incredibly attractive if they hold in real trading. As we approach an era where these prototypes go live, a key watchpoint will be whether they maintain these stats out-of-sample.

## Academic Prototypes vs. Real-World Deployments

There is a noticeable gap between **academic research prototypes** and **deployed real-world AI investor systems**, though it's narrowing rapidly. Understanding their differences can help in designing a personal AI investor:

- **Data and Environment:** Academic systems often backtest on historical data and sometimes simulate ideal conditions (low transaction costs, ample liquidity). Real deployments face slippage, market impact, and occasionally data quality issues. For instance, an RL agent that performed well on minute-by-minute stock data in a lab might struggle with broker execution delays or unexpected news spikes. Hedge funds tackle this by incorporating execution cost models and sometimes **training agents specifically for execution tasks** (e.g. RL for optimal order slicing). The FinRL Contest series emphasizes realistic market environments and even had an “**order execution challenge**” to address this <sup>16</sup>. In building your own AI investor, you'd need to include modules for order execution or use a broker API with safeguards (like limit orders, etc.). Open-source projects such as *ElegantRL* and *FinRL* provide simulated brokerage environments to test these aspects.
- **Architectural Complexity:** Academic multi-agent systems can be quite complex (dozens of agents, multiple layers of communication), which is feasible in a controlled experiment. Real hedge funds often start simpler and only add complexity if it pays off. Many successful quant funds use ensembles of relatively simple models – effectively a multi-agent approach, but not necessarily with agents chatting in natural language. There's a trade-off between **transparency** and **performance**: the multi-LLM systems (like TradingAgents) prioritize human-like reasoning and explainability, which is appealing but computationally heavy. In contrast, a fund like Renaissance Technologies famously uses large ensembles of machine learning models whose internal logic is opaque, but they care only about the predictive power. However, with increasing regulatory focus on AI explainability in finance <sup>76</sup> <sup>75</sup>, the academic trend toward explainable agents might influence industry practice. A personal AI investor could strike a middle ground – e.g. use a handful of specialized agents you can monitor, rather than an inscrutable black-box.
- **Resource Constraints:** Real-world systems must consider runtime speed and cost. An academic might run an expensive deep learning model on historical data without issue, but a live system

needs to operate within cost budgets (especially if using paid LLM APIs or requiring GPUs). The TradingAgents paper noted issues like *latency* and *cost of inference at scale* as challenges for multi-agent LLM trading bots <sup>58</sup>. Solutions include optimizing prompts, using smaller fine-tuned models for specific tasks (e.g. a lightweight sentiment model instead of a giant LLM), and triggering agents only when needed (event-driven activation). A DIY AI investor can be designed to run mostly on your local machine or a cloud VM, updating perhaps a few times a day, which keeps costs low while still being effective for swing trading or daily rotations.

- **Validation and Adaptation:** Academic results are typically validated on past data. In the real world, continuous learning or adaptation is crucial. Hedge funds might retrain models periodically or have self-learning pipelines. A personal AI investor should be set up with a **training-testing-trading pipeline** that allows updating the model as new data arrives while avoiding overfitting. The FinRL framework explicitly promotes this pipeline: train on one period, test on another, then deploy and monitor <sup>77</sup> <sup>78</sup>. It also encourages **modularity** and “plug-and-play” experimentation – e.g. you can easily swap out the RL algorithm or add a new data source without rebuilding everything <sup>79</sup> <sup>80</sup>. In practice, many open-source projects (e.g. *Freqtrade* for crypto) allow you to backtest and forward-test strategies before going live. Emulating that rigor from academia in your personal system is wise: maintain a clear distinction between in-sample training data and out-of-sample evaluation, and use paper trading to validate performance in real market conditions.
- **Examples of Real Deployments:** Beyond hedge funds, there are fintech products like **AI-enhanced ETFs**. For instance, Qraft’s AI-Driven ETFs use machine learning to pick stocks and have in some periods beaten the S&P 500 (though not consistently every year). There are also decentralized experiments: protocols that let users allocate to AI-managed crypto portfolios (some on Numerai or token sets). These often leverage simpler strategies or crowd-sourced models, not full multi-agent AI. One noteworthy concept is **Decentralized Autonomous Funds (DAFs)** where a pool of capital is managed by on-chain AI agents with minimal human oversight <sup>81</sup> <sup>82</sup>. While still experimental, it hints that “hands-off” AI investing is being pursued in practice. The Kenson Investments whitepaper (2025) suggests that **blockchain and smart contracts** might be integrated with AI to create fully autonomous funds – essentially code that invests money according to AI signals, with transparency and immutability on-chain <sup>76</sup> <sup>82</sup>.

Overall, the trajectory is that academic prototypes introduce innovative ideas (like multi-agent debates, hierarchical RL, etc.), and then the best of those get simplified or optimized into real systems. Hedge funds often operate in secrecy (“money grows in the dark” as one expert said <sup>83</sup>), so we only hear of successes or failures anecdotally. But a private individual today has access to many of the same tools via open source: for example, **FinRL** library (by Columbia University) provides ready-to-use DRL algorithms and market gym environments; **LangChain/LangGPT** frameworks allow building LLM-driven agents with relative ease; cloud platforms (like DigitalOcean’s Gradient, mentioned in the TradingAgents article) offer infrastructure for deploying custom AI agents without much DevOps overhead <sup>84</sup> <sup>85</sup>. The democratization of these technologies means a motivated individual can prototype their own AI investor akin to a small-scale hedge fund AI.

## Building a Personal AI Investor – Key Takeaways

Designing a personal AI investor system on one's own infrastructure is an ambitious but increasingly feasible project. Based on the review above, here are key considerations and recommended components for such a system:

- **Modular Multi-Agent Architecture:** Structure your system into modules that reflect the trading process. For example, you might have: **Data Ingestion agents** (one for market price data, one for news sentiment, one for alternative data like on-chain stats), a **Strategy agent** (or multiple strategy agents) that proposes trades based on the processed inputs, a **Risk Management agent** that evaluates proposals against risk limits (e.g. max drawdown threshold, position size limits), and an **Execution module** that places trades via your broker or exchange API. This modular design is aligned with both research and industry practices <sup>20</sup> <sup>86</sup>. Each agent can be developed and improved independently, and you can add new agents over time (for example, if you want to incorporate options trading later, you can bolt on an "Options strategist" agent). Keeping agents separate also means you can use different methods for each: perhaps a small neural network for technical signals, an API call to a sentiment service for news, and a rule-based script for risk checks – all orchestrated by a top-level program.
- **Use of Reinforcement Learning and Supervised Models:** For the strategy agent(s), leveraging DRL is a natural choice if you have sufficient data to train on. You could use libraries like FinRL or Stable Baselines3 to train an agent in a **market environment** that represents the assets you care about (there are open data and pre-built environments for stocks, crypto, etc.) <sup>87</sup> <sup>77</sup>. Ensure you define the reward with your goals in mind – e.g. include a penalty for large drawdowns or a term for Sharpe ratio to guide the agent toward the risk/return profile you want. Alternately or additionally, incorporate some **supervised learning models**: for example, a price predictor for the next day (using an LSTM or transformer on historical data) whose output is one input into your trading decision. Many successful systems are ensembles, so you might combine an RL policy's action with signals from a few predictive models (this could be as simple as: take the RL agent's trade only if confirmed by a trend predictor). This was seen in some contest solutions where participants blended methods for more robustness <sup>88</sup> <sup>50</sup>.
- **LLM for Information and Explanations:** Including an LLM component can greatly enhance the system's scope. You can use an LLM (like GPT-4 via API, or a smaller local model for cost savings) to **digest news headlines daily** and output a summary sentiment or list of concerns/opportunities. This can be fed into the strategy agent (e.g. as an extra observation feature indicating "news sentiment = positive" or "big geopolitical risk event = true"). You can also have the LLM serve as a **commentator** on the model's decisions – after your system decides on trades for the day, prompt the LLM to explain those trades in plain language ("The AI is buying X because..."). This not only helps you trust the system but also flags if something is off (if the explanation is nonsense, maybe the decision is suspect). P1GPT demonstrated how having explanations can make an AI trader more transparent and potentially more trustworthy for users <sup>53</sup> <sup>47</sup>. Just be mindful to constrain the LLM's role to what it's good at (text and high-level reasoning) and not let it directly control money without additional checks, as LLMs can sometimes produce inconsistent outputs if prompts are ambiguous.

- **Backtesting and Evaluation:** Before live deployment, rigorously backtest your system on historical data and *out-of-sample* periods. You might segment data into a training period for the models (e.g. 2015–2020) and a testing period (2021–2023) to see how it would have done more recently. Use standard metrics: annualized return, volatility, Sharpe ratio, maximum drawdown, Calmar ratio, etc. Aim to meet or exceed that **2x SPY return, <20% drawdown** in backtests first, with a decent margin for error (because real trading tends to perform a bit worse than backtest due to frictions). If possible, do **paper trading** (live simulation) for a few weeks to validate that the strategy behaves as expected with live data and order execution. Many DIY traders use platforms like Interactive Brokers' paper trading or crypto exchange testnets for this purpose.
- **Risk Management and Human Oversight:** Even though the goal is an automated system, it's wise to have some "**circuit breakers**" and alerts. For example, program the system to halt trading or notify you if the portfolio drops more than e.g. 5% in a day or if the AI outputs a very unusual trade (like suddenly allocating 100% to a single penny stock – which a bug or outlier input might cause). This aligns with professional practice: AI funds often have human risk managers overseeing the AI's positions, ready to intervene if needed <sup>89</sup>. You can keep the intervention minimal (the ideal is once or twice a day check-in, as you mentioned), but having a dashboard to monitor positions, P/L, and agent reasoning can be invaluable. The modular design helps here: you can log each agent's output (e.g. "Technical agent says buy, News agent sentiment = very negative, Risk agent veto = false") to understand how a decision was made.
- **Infrastructure:** Running a multi-agent AI on your own infrastructure is doable on a modern PC or modest cloud server for low-frequency trading. You'll need a data pipeline (for equities, APIs like Alpha Vantage or Yahoo Finance can fetch daily/hourly data; for crypto, exchange APIs or aggregators like CoinGecko; for news, RSS feeds or news APIs; for social sentiment, perhaps Twitter API or alternative datasets). A database or at least in-memory data store will be needed to feed your agents. The FinRL documentation emphasizes a **three-layer architecture** (Data layer -> Environment layer -> Agent layer) which is a good template <sup>90</sup>. In practice: your Data layer fetches and updates datasets, the Environment layer formulates the state (calculating indicators, etc.), and the Agent layer makes decisions. This separation makes the system easier to maintain. For instance, if you switch brokers, you only update the Execution part without touching how the agent decides. If you add a new data source, you plug it into the Data layer and retrain or re-run the agent. Consider using scheduling (e.g. a daily cron job) to run the process if it's not continuous, or event-driven triggers if you want it to react to certain events (like a news alert).
- **Continuous Learning vs. Static Model:** Decide if your AI will retrain periodically on new data. Some personal projects simply train a model once and then trade with it until performance degrades. More sophisticated setups retrain maybe monthly or when triggered by significant changes. Reinforcement learning agents can be set to *online learn* (updating after each trade) but that's risky without careful safeguards (it could learn from its own mistakes incorrectly). A safer approach is **incremental retraining**: e.g. every week, incorporate the latest week of data, retrain the model (or fine-tune it) and update the agent if it improves validation metrics. With modern cloud compute, you could even automate this. The FinRL-Meta initiative moves in this direction by automating data updates and model re-training in a pipeline fashion <sup>91</sup> <sup>92</sup>.

In conclusion, building a personal AI investor draws on the same principles fueling AI-driven hedge funds and research prototypes: modular design, multi-disciplinary intelligence (combining technical signals with

fundamental and sentiment analysis), and rigorous training and risk control. The exciting news is that the **tools and knowledge are accessible** – one can leverage open-source libraries for RL, pre-trained financial LLMs like FinGPT for language analysis, and cloud platforms for deployment. By studying successful examples (academic papers and the first generation of AI funds), an individual can assemble a “homebrew” AI investor that, with some luck and skill, approaches the performance of much larger institutional systems. As always, caution is warranted – extensive testing and a clear understanding of your AI’s failure modes are key. But the path to an largely **hands-off AI portfolio manager** beating the market is increasingly within reach, powered by the advances in deep learning, multi-agent coordination, and NLP integration that we’ve discussed. The **future of investing may well belong to such autonomous AI agents**, and building one yourself is a step into that future.

## References

- Liu, J. (2025). *Multi-Agent AI Architecture for Personalized DeFi Investment Strategies*. (SuperIntent case study) 23 10
- Intellectyx AI. (2025). *TradingAgents: A Multi-Agent LLM Financial Trading Framework*. (Medium article) 93 9
- Khan, F. (2025). *Building a Deep Thinking Trading System with Multi-Agentic Architecture*. (Concept demo) 6 7
- Lu, C.C., et al. (2025). *P1GPT: A Multi-Agent LLM Workflow for Multi-Modal Financial Analysis*. (arXiv 2510.23032) 54 30
- Kim, S.H. & Lee, K.H. (2025). *MAMA: Multi-Asset Multi-Agent Reinforcement Learning for Portfolio Management*. (IEEE Access) 31
- Tan, B., et al. (2024). *OPHR: Mastering Volatility Trading with Multi-Agent Deep RL*. (Preprint) 24 28
- Wang, K., et al. (2025). *FinRL Contests: Benchmarking Financial RL Agents 2023–2025*. 16 33
- Jess L. (2025). *Your Guide to the TradingAgents Multi-Agent LLM Framework*. (DigitalOcean blog) 50 5
- Kenson Investments. (2025). *The AI-Powered Hedge Fund: How ML is Reshaping Investment Strategies*. (Whitepaper) 74 75
- Georgia M. (2015). *Artificial intelligence is the next big thing for hedge funds...* (Quartz) 64 68

---

1 2 3 9 11 58 70 71 93 TradingAgents: A Multi-Agent LLM Financial Trading Framework | by Intellectyx AI | Dec, 2025 | Medium

<https://medium.com/@intellectyxai/tradingagents-a-multi-agent-lm-financial-trading-framework-78d08acfef63>

4 5 32 44 49 50 51 73 84 85 Your Guide to the TradingAgents Multi-Agent LLM Framework | DigitalOcean

<https://www.digitalocean.com/resources/articles/tradingagents-lm-framework>

6 7 72 Building a Deep Thinking Trading System with Multi-Agentic Architecture | by Fareed Khan | Level Up Coding

<https://levelup.gitconnected.com/building-a-deep-thinking-trading-system-with-multi-agentic-architecture-c13da7effd2d?gi=19e5c9d016fc>

8 10 17 18 20 21 22 23 55 56 86 Multi-Agent AI Architecture for Personalized DeFi Investment Strategies | by Jung-Hua Liu | Medium

<https://medium.com/@gwxr2005/multi-agent-ai-architecture-for-personalized-defi-investment-strategies-c81c1b9de20c>

- 12 13 14 15 29 30 42 45 52 53 54 57 59 60 61 62 P1GPT: A Multi-Agent LLM Workflow Module for Multi-Modal Financial Information Analysis  
<https://arxiv.org/html/2510.23032v1>
- 16 33 43 88 FinRL Contests: Benchmarking Data-driven Financial Reinforcement Learning Agents  
<https://arxiv.org/html/2504.02281v3>
- 19 40 A multi-agent virtual market model for generalization in ...  
<https://www.sciencedirect.com/science/article/abs/pii/S1568494623000030>
- 24 25 26 27 28 63 openreview.net  
<https://openreview.net/pdf?id=2p4AtivyZz>
- 31 38 39 (PDF) Multi-Asset Multi-Agent Reinforcement Learning for Portfolio Management  
[https://www.researchgate.net/publication/397593372\\_Multi-Asset\\_Multi-Agent\\_Reinforcement\\_Learning\\_for\\_Portfolio\\_Management](https://www.researchgate.net/publication/397593372_Multi-Asset_Multi-Agent_Reinforcement_Learning_for_Portfolio_Management)
- 34 FinRL's Implementation of DRL Algorithms for Stock Trading  
<https://hackernoon.com/finrls-implementation-of-drl-algorithms-for-stock-trading>
- 35 64 65 66 67 68 69 83 89 Artificial intelligence is the next big thing for hedge funds seeking an edge  
<https://qz.com/389647/artificial-intelligence-is-the-next-big-thing-for-hedge-funds-seeking-an-edge>
- 36 Research articles | Scientific Reports - Nature  
<https://www.nature.com/srep/research-articles/1000?searchType=journalSearch&sort=PubDate&type=article&year=2026&page=14>
- 37 Articles in 2026 | Scientific Reports - Nature  
<https://www.nature.com/srep/articles?searchType=journalSearch&sort=PubDate&year=2026&page=9>
- 41 Multi-Agent Reinforcement Learning for Market Making - arXiv  
<https://arxiv.org/html/2510.25929v1>
- 46 47 P1GPT: a multi-agent LLM workflow module for multi-modal financial information analysis  
<https://arxiv.org/pdf/2510.23032>
- 48 [PDF] TradingAgents: Multi-Agents LLM Financial Trading Framework - arXiv  
<https://arxiv.org/pdf/2412.20138>
- 74 75 76 81 82 The AI-Powered Hedge Fund: How Machine Learning is Reshaping Investment Strategies - kenson Investments  
<https://kensoninvestments.com/resources/the-ai-powered-hedge-fund-how-machine-learning-is-reshaping-investment-strategies/>
- 77 78 79 80 87 90 91 92 Overview — FinRL 0.3.1 documentation  
[https://finrl.readthedocs.io/en/latest/finrl\\_meta/overview.html](https://finrl.readthedocs.io/en/latest/finrl_meta/overview.html)